

Introduction to AgentPack

AgentPack is a document-to-agent-context compiler. It ingests PDFs, Markdown, CSV, and plain text files, chunks them into semantically coherent blocks, and produces a manifest that LLM agents can query.

Architecture Overview

The pipeline consists of four stages: parsing, chunking, embedding, and indexing. Each stage is independently testable and produces well-typed intermediate artefacts.

Supported Formats

Format	Parser	Semantic
PDF	PDFParser	Yes (Docling)
Markdown	MarkdownParser	Yes
CSV	CSVParser	Table
TXT	TextParser	No

Conclusion

AgentPack is designed to be extended. New parsers register themselves via a simple interface. The retrieval layer supports both lexical and semantic search.